

Distributed FFTs for integrated algorithms

Tze Meng Low
Elliott Binder




Planewave Pseudopotential Method in DFT

$$\left\{ -\frac{1}{2} \nabla^2 + \int \frac{\rho(r')}{|r-r'|} dr' + \sum_I \frac{Z}{|r-R_I|} + V_{xc}(\rho(r)) \right\} \psi_j(r) = E_j \psi_j(r)$$

<u>Computational Task (CG solver)</u>	<u>Scaling</u>
Orthogonalization	MN^2 ($\propto N^3$)
Subspace diagonalization	N^3
3d FFTs (most communications)	$NM \log M$ ($\propto N^2$)
Nonlocal pseudopotential	MN^2 (N^2 real space)

N: number of eigenpairs (bands, states, electrons) required (lowest in spectrum) (N is in the range of 100s to 1000s)

M: matrix (Hamiltonian) dimension, #Fourier components ($M \sim 256^3$ to 1024^3)



Basic 3DFFT

- 3 functions to invoke distributed 3DFFT. Similar to all many distributed 3DFFT framework


```
fftx_plan plan = fftx_plan_distributed(r, c, M, N, K);
```

```
fftx_execute(plan, (double*)device_out_buffer,  
             (double*)device_in_buffer, FFTX_FORWARD);
```

```
fftx_plan_destroy(plan);
```

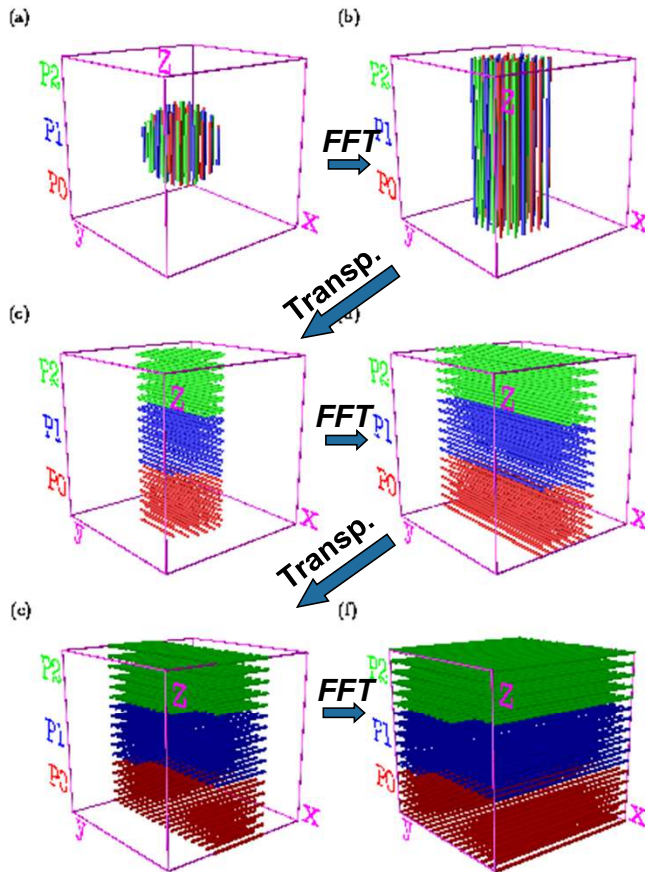
- Assumptions

- Data is evenly distributed in blocks of pencils
- All temporary data and communication buffers will be created as part of the plan
- Requires at least $2 \times (MNK / (rc))$ additional space on the GPU for intermediate results



Requirements for Integrated Planewave algorithm

FIGURES



- Many zero elements in data cube
- Exploiting zero elements reduces computation and communication

Communications

- Ratio of Cube to Cylinder Volume is $\pi/16 = 0.196$
- **5.1 factor less** data communicated compared to full (library) 3D FFT

FLOPS

- Three sets of 1D FFTS $(\pi/16)N^2, (1/2)N^2, N^2$
- Full 3D FFT $3N^2$ sets of 1D FFTS
- **1.70 factor lower** FLOP count compared to full (library) 3D FFT

Embedded 3DFFT


- Similar interface with `is_embedded` set to `true`

```
fftx_plan plan = fftx_plan_distributed(r, c, M, N, K, is_embedded);
```

```
fftx_execute(plan, (double*)device_out_buffer,  
             (double*)device_in_buffer, FFTX_FORWARD);
```

```
fftx_plan_destroy(plan);
```

- Exploit locations of non-zeros to reduce amount of data communicated
- Padding performed for 2nd and 3rd stage performed as part of the communication
- Same assumptions as the basic 3DFFT



Backup Slides with code

```
fftx_plan plan = fftx_plan_distributed(r, c, M, N, K, batch, is_embedded);

DEVICE_SYNCHRONIZE();
MPI_Barrier(MPI_COMM_WORLD);

for (int t = 0; t < 1; t++) {

    double start_time = MPI_Wtime();

    fftx_execute(plan, (double*)out_buffer, (double*)in_buffer, (is_forward ? FFTX_FORWARD: FFTX_BACKWARD));

    double end_time = MPI_Wtime();

    double min_time    = min_diff(start_time, end_time, MPI_COMM_WORLD);
    double max_time    = max_diff(start_time, end_time, MPI_COMM_WORLD);

    DEVICE_MEM_COPY(fftx_out, out_buffer, (Mo*No*Ko/p) * sizeof(complex<double>)*batch, MEM_COPY_DEVICE_TO_HOST);
    DEVICE_SYNCHRONIZE();

}
```

Execute with command

```
srun -A CSC304_crusher -N 4 -n $(( $P * $P )) --gpus-per-node 8 --gpu-bind=closest -t 00:01:00 ./testmddft3d.x
```

Performance on Crusher

Full vs Embedded 3DFFT

