

# FFTX Backend: SPIRAL

Presenter: Franz Franchetti



# Have You Ever Wondered About This?

## Numerical Linear Algebra

**LAPACK**  
**ScaLAPACK**  
LU factorization  
Eigensolves  
SVD

**BLAS, BLACS**  
BLAS-1  
BLAS-2  
BLAS-3

## Spectral Algorithms

Convolution  
Correlation  
Upsampling  
Poisson solver  
...



**FFTW**  
DFT, RDFT  
1D, 2D, 3D,...  
batch



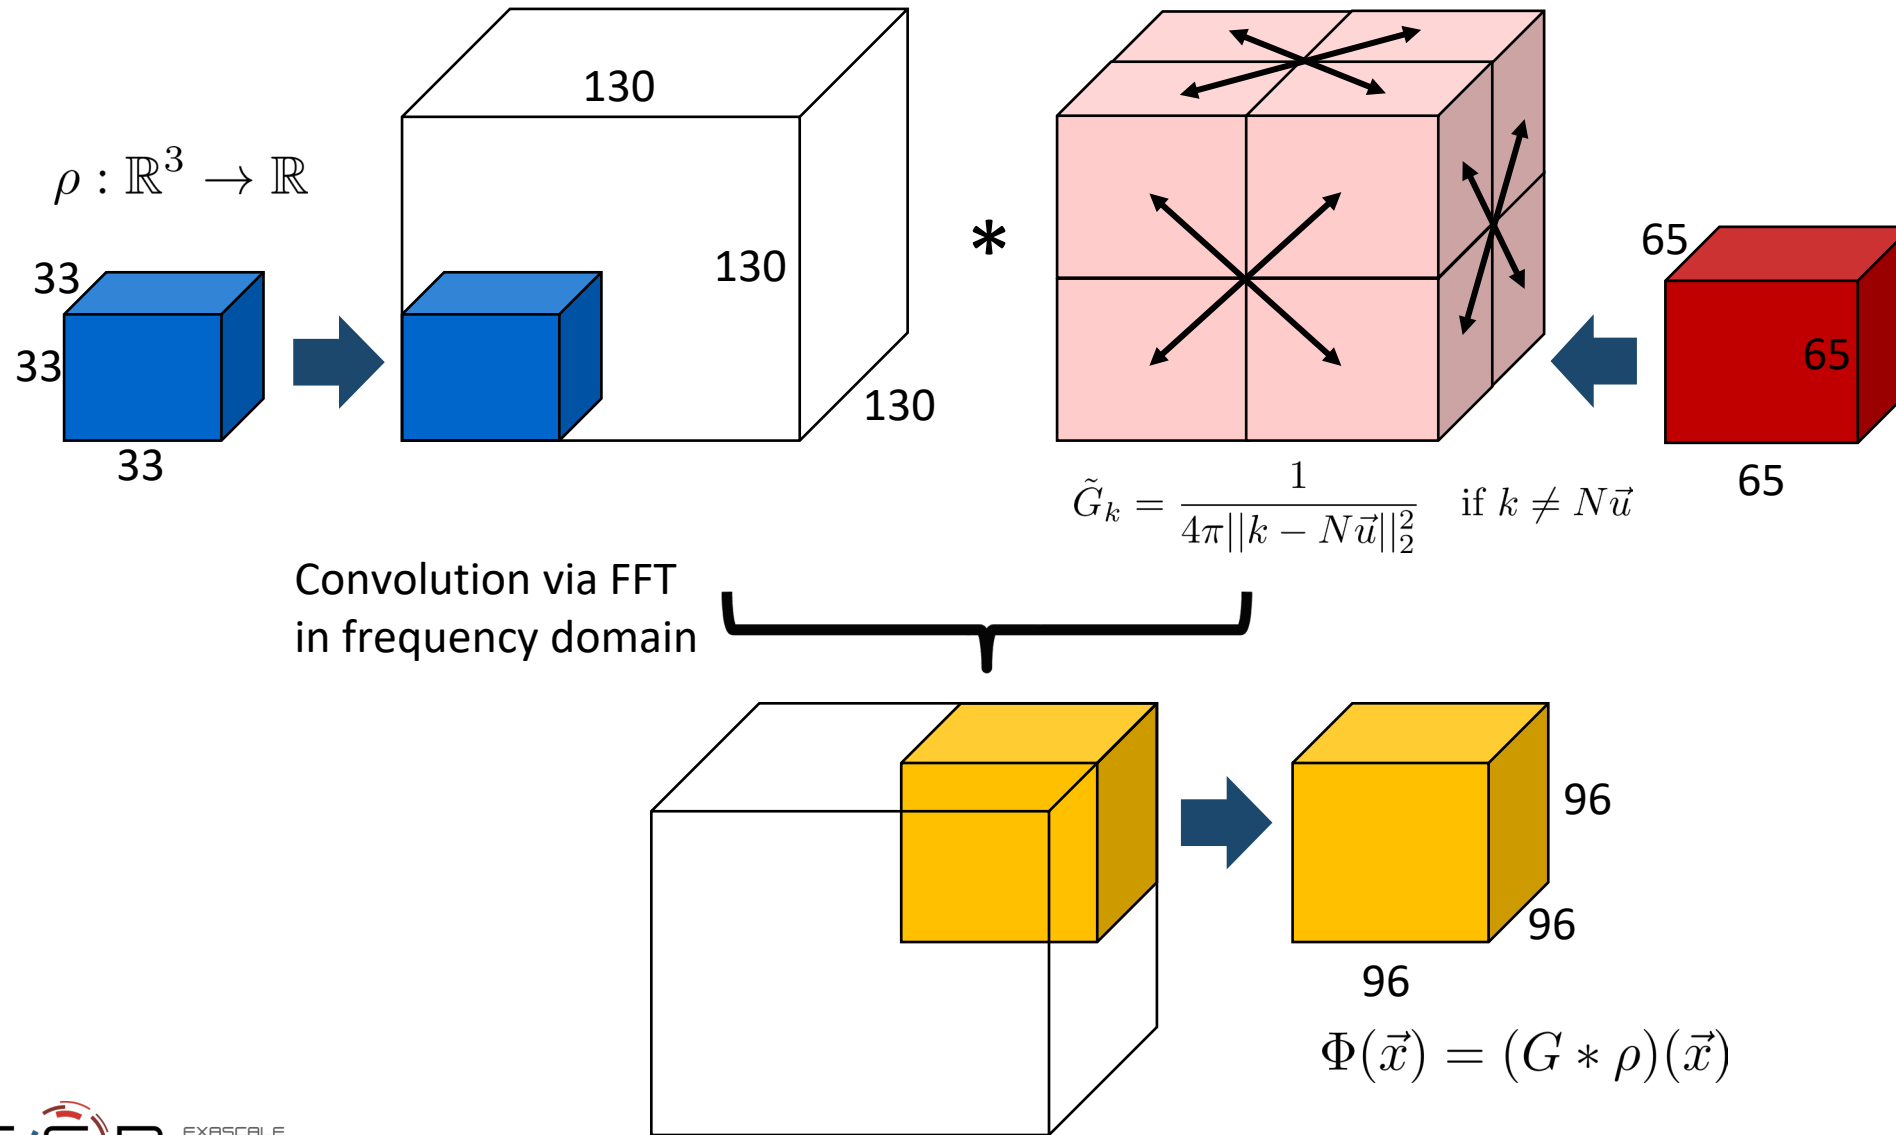
**SpectralPACK**  
Convolution  
Correlation  
Upsampling  
Poisson solver  
...

**FFTX**  
DFT, RDFT  
1D, 2D, 3D,...  
batch

## No LAPACK equivalent for spectral methods

- Medium size 1D FFT (1k—10k data points) is most common library call  
applications break down 3D problems themselves and then call the 1D FFT library
- Higher level FFT calls rarely used  
FFTW *guru* interface is powerful but hard to used, leading to performance loss
- **Low arithmetic intensity and variation of FFT use make library approach hard**  
Algorithm specific decompositions and FFT calls intertwined with non-FFT code

# Algorithm: Hockney Free Space Convolution



# Defining / Generating / Using an FFTX Transform

Defining the DAG:

```
const int nx=80; const int ny=80; const int nz=80;
box_t<3> domain(point_t<3>({{1,1,1}}), point_t<3>({{nx,ny,nz}}));
array_t<3,std::complex<double>> inputs(domain);
array_t<3,std::complex<double>> outputs(domain);
std::array<array_t<3,std::complex<double>>,2> intermediates {domain};
setInputs(inputs); setOutputs(outputs); openScalarDAG();
MDDFT(domain.extents(), 1, intermediates[0], inputs);
RCDiag(domain.extents(), 1, intermediates[1], intermediates[0]);
IMDDFT(domain.extents(), 1, outputs, intermediates[1]);
closeScalarDAG(intermediates, "mdconv");
```

Generated code is accessed  
by an application through a  
C++ class API.

```
fftx::point_t<3> sz3d = {nx,ny,nz};
fftx::box_t<3> domain(fftx::point_t<3>({{1,1,1}}), sz3d);

fftx::array_t<3,std::complex<double>> input(domain);
fftx::array_t<3,std::complex<double>> output(domain);
fftx::array_t<3,std::complex<double>> filter(domain);

fftx::mdconv<3> mdconv3d(sz3d); // library API.
mdconv3d.transform(input,output,filter); // Call convolution.
```

FFTX / Spiral workflow generates  
 $O(10^3)$  lines per size of C++/HIP/cuda/SYCL

```
t1629 = (t1609 + t1614);
t1630 = (s854 + s864);
t1631 = (s855 - s865);
t1632 = (s854 - s864);
t1633 = (s855 + s865);
Y[(a2899 + 256)] = (t1626 + t1630);
Y[(a2899 + 257)] = (t1627 + t1631);
Y[(a2899 + 1280)] = (t1626 - t1630);
Y[(a2899 + 1281)] = (t1627 - t1631);
Y[(a2899 + 768)] = (t1628 + t1633);
Y[(a2899 + 769)] = (t1629 - t1632);
Y[(a2899 + 1792)] = (t1628 - t1633);
Y[(a2899 + 1793)] = (t1629 + t1632);
}
}

void mddft3d_1024x1024x1024_mpi(double *Y, double *X) {
    dim3 b458(1, 1, 64), b459(1, 1, 1), b460(1, 1, 64), b461(1, 1, 1),
    1),
    g4(1, 1, 1), g5(64, 256, 1);
    ker_code0<<<g1, b458>>>(X);
    fftx_mpi_rcperm(Q2, Q1, 1073741824, 2, 3, 1024, 1024, 1024);
    ker_code2<<<g3, b460>>>();
    fftx_mpi_rcperm(Q2, Q1, 1073741824, 1, 3, 1024, 1024, 1024);
    ker_code4<<<g5, b462>>>(Y);
```

# FFTX C++ Code: Defining a FFTX Transform

```
#include "fftx3.hpp"
...
int main(int argc, char* argv[])
{
    tracing=true;
    const int nx=80; const int ny
    box_t<3> domain(point_t<3>({{

    array_t<3,std::complex<double
    array_t<3,std::complex<double
    std::array<array_t<3,std::com

    setInputs(inputs);
    setOutputs(outputs);
    openScalarDAG();

    MDFFT(domain.extents(), 1, intermediates[0], inputs);
    RCDiag(domain.extents(), 1, intermediates[1], intermediates[0]);
    IMDDFT(domain.extents(), 1, outputs, intermediates[1]);

    closeScalarDAG(intermediates, "mdconv");
```

## This is a specification dressed as a program

- Needs to be clean and concise
- No code level optimizations and tricks
- Don't think "performance" but "correctness"
- *For production code and software engineering*

# C/C++ FFTX Program Trace: MPI+GPU

```
Load(fftx);  
Load(mpi);  
Import(fftx, mpi);
```

```
pg := [32,32];  
d := 3;  
N := Replicate(d, 1024);  
procGrid := MPIProcGrid
```

```
name := "mdconv1024x1024";  
conf := LocalConfig.mp
```

```
Nlocal := N{[1]}::List;  
localBrick := TArrayND;  
dataLayout := TPtrGlob  
Xglobal := tcast(dataL  
Yglobal := tcast(dataL
```

```
Tlglobal := var("T1", dataLayout); T2global := var("T2", dataLayout);  
symvar := var("symvar", dataLayout);  
t := TFCall(TDecl(TDAG([  
  TDAGNode(TRC(MDDFT(N, -1)), Tlglobal, Xglobal),  
  TDAGNode(RCDiag(FDataOfs(tcast(TPtr(TReal), symvar), 2*Product(N), 0)), T2global, Tlglobal),  
  TDAGNode(TRC(MDDFT(N, 1)), Yglobal, T2global)  
]), [ Tlglobal, T2global ]), rec(fname := name, params := [ symvar ]));
```

```
opts := conf.getOpts(t);  
tt := opts.tagIt(t);  
c := opts.fftxGen(tt);  
opts.prettyPrint(c);  
PrintTo(name:".cu", opts.prettyPrint(c));
```

## The whole convolution kernel is captured

- DAG with all dependencies
- User-defined call-backs
- Captures pruning, zero-padding and symmetries
- *Lifts sequence of C/C++ library calls to a specification*

**Operator DAG abstraction:**

**Major new concept and capability in SPIRAL due to FFTX**

# Advanced Operator DAG: WarpX Convolution

```

TFCall (TDecl (TDAG ( [
  TDAGNode (TResample([n, n, n], [np, np, n], [0.0, 0.0, -0.5]), nth(boxBig0, 0), nth(X, 0)),
  TDAGNode (TResample([n, n, n], [np, n, np], [0.0, -0.5, 0.0]), nth(boxBig0, 1), nth(X, 1)),
  TDAGNode (TResample([n, n, n], [n, np, np], [-0.5, 0.0, 0.0]), nth(boxBig0, 2), nth(X, 2)),
  TDAGNode (TResample([n, n, n], [n, n, np], [-0.5, -0.5, 0.0]), nth(boxBig0, 3), nth(X, 3)),
  TDAGNode (TResample([n, n, n], [n, np, n], [-0.5, 0.0, -0.5]), nth(boxBig0, 4), nth(X, 4)),
  TDAGNode (TResample([n, n, n], [np, n, n], [0.0, -0.5, -0.5]), nth(boxBig0, 5), nth(X, 5)),
  TDAGNode (TResample([n, n, n], [np, np, n], [0.0, 0.0, -0.5]), nth(boxBig0, 6), nth(X, 6)),
  TDAGNode (TResample([n, n, n], [np, n, np], [0.0, -0.5, 0.0]), nth(boxBig0, 7), nth(X, 7)),
  TDAGNode (TResample([n, n, n], [n, np, np], [-0.5, 0.0, 0.0]), nth(boxBig0, 8), nth(X, 8)),
  TDAGNode (TResample([n, n, n], [np, np, np], [0.0, 0.0, 0.0]), nth(boxBig0, 9), nth(X, 9)),
  TDAGNode (TResample([n, n, n], [np, np, np], [0.0, 0.0, 0.0]), nth(boxBig0, 10), nth(X, 10)),
  TDAGNode (TTensorI(MDPRDFT([n, n, n], -1), inFields, APar, APar), boxBig1, boxBig0),
  TDAGNode (TRC(TMap(rmat, [iz, iy, ix], AVec, AVec)), boxBig2, boxBig1),
  TDAGNode (TTensorI(IMPDRDFT([n, n, n], 1), outFields, APar, APar), boxBig3, boxBig2),
  TDAGNode (TResample([np, np, n], [n, n, n], [0.0, 0.0, 0.5]), nth(Y, 0), nth(boxBig3, 0)),
  TDAGNode (TResample([np, n, np], [n, n, n], [0.0, 0.5, 0.0]), nth(Y, 1), nth(boxBig3, 1)),
  TDAGNode (TResample([n, np, np], [n, n, n], [0.5, 0.0, 0.0]), nth(Y, 2), nth(boxBig3, 2)),
  TDAGNode (TResample([n, n, np], [n, n, n], [0.5, 0.5, 0.0]), nth(Y, 3), nth(boxBig3, 3)),
  TDAGNode (TResample([n, np, n], [n, n, n], [0.5, 0.0, 0.5]), nth(Y, 4), nth(boxBig3, 4)),
  TDAGNode (TResample([np, n, n], [n, n, n], [0.0, 0.5, 0.5]), nth(Y, 5), nth(boxBig3, 5)) ]),
[boxBig0, boxBig1, boxBig2, boxBig3]),
rec(XType := TPtr(TPtr(TReal)), YType := TPtr(TPtr(TReal)),
  fname := name, params := [symvar ])

```



# Advanced DAG: WarpX Pointwise Kernel

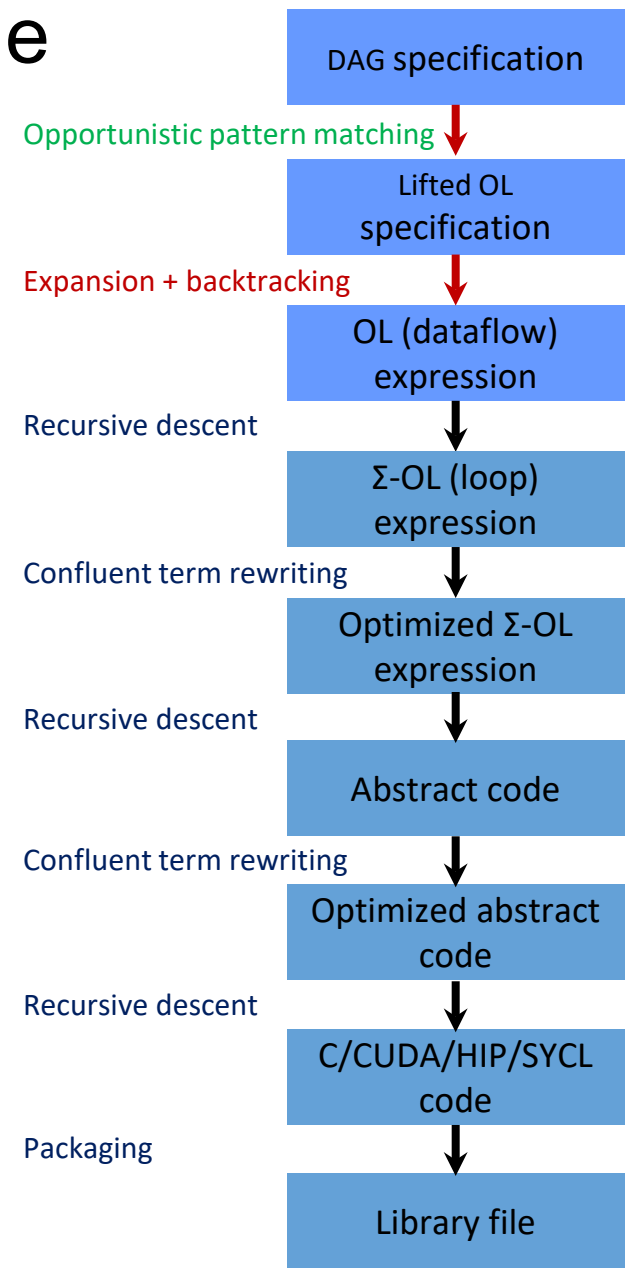
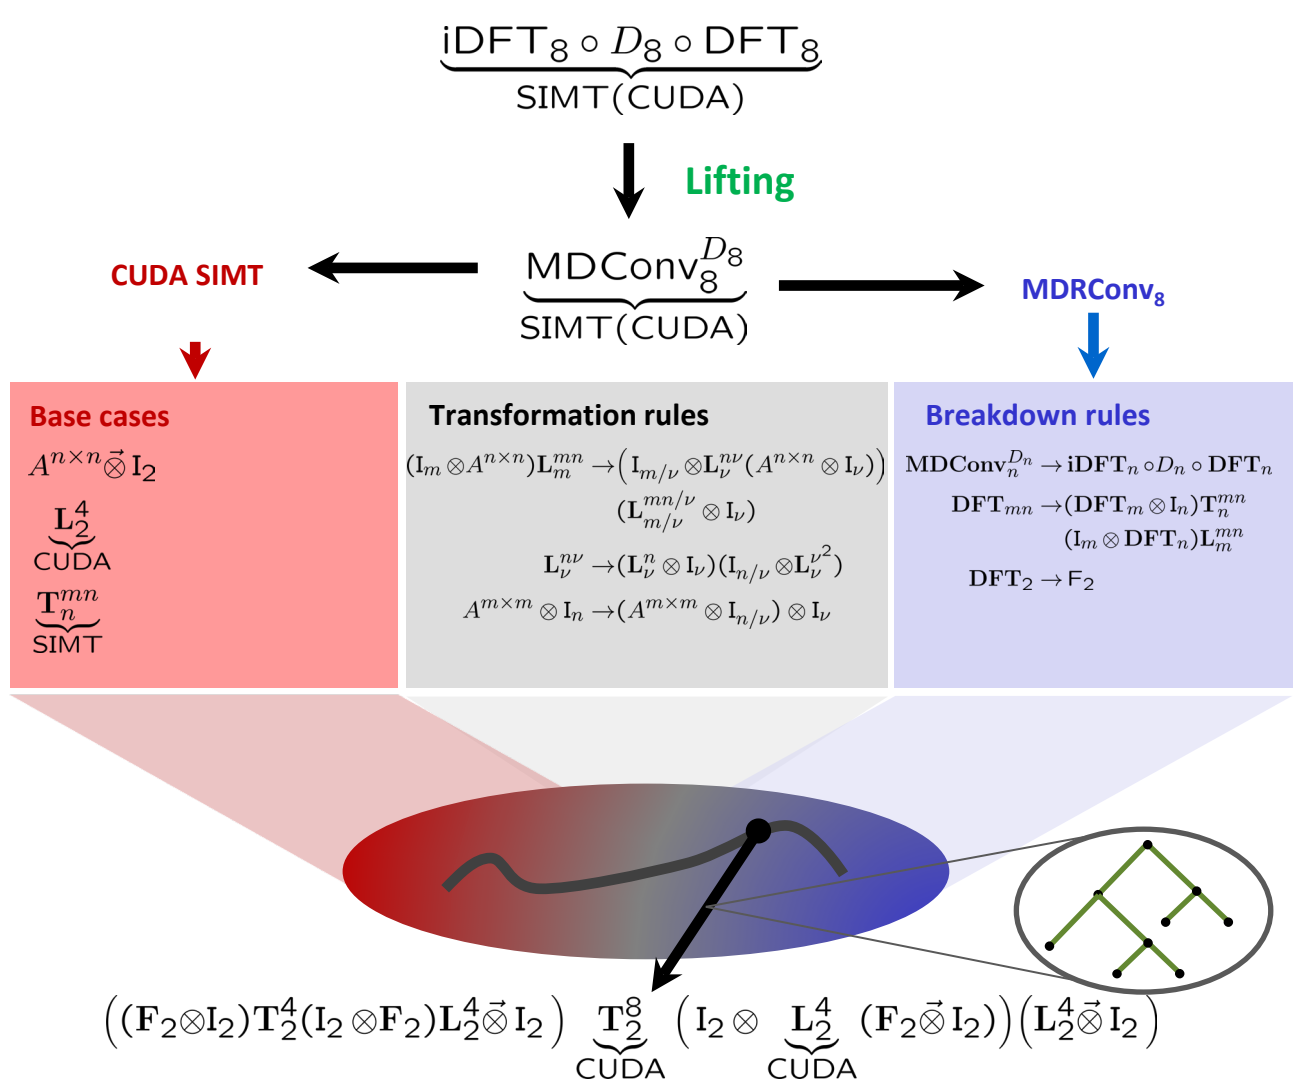
```
ix := Ind(xdim), iy := Ind(ydim), iz := Ind(zdim), ii := lin_idx(iz, iy, ix),
fmkx := nth(nth(symvar, 0), ix), fmky := nth(nth(symvar, 1), iy),
fmkz := nth(nth(symvar, 2), iz), fcv := nth(nth(symvar, 3), ii),
fsckv := nth(nth(symvar, 4), ii), fx1v := nth(nth(symvar, 5), ii),
fx2v := nth(nth(symvar, 6), ii), fx3v := nth(nth(symvar, 7), ii),
```

```
rmat := TSparseMat([outFields,inFields], [
[0, [0, fcv / n^3], [4, cpack(0, -fmkz * c2 * fsckv / n^3)],
[5, cpack(0, fmky * c2 * fsckv / n^3)], [6, -invep0 * fsckv / n^3],
[9, cpack(0, fmkx * fx3v / n^3)], [10, cpack(0, -fmkx * fx2v / n^3)]],
[1, [1, fcv / n^3], [3, cpack(0, fmkz * c2 * fsckv / n^3)],
[5, cpack(0, -fmkx * c2 * fsckv / n^3)], [7, -invep0 * fsckv / n^3],
[9, cpack(0, fmky * fx3v / n^3)], [10, cpack(0, -fmky * fx2v / n^3)]],
[2, [2, fcv / n^3], [3, cpack(0, -fmky * c2 * fsckv / n^3)],
[4, cpack(0, fmkx * c2 * fsckv / n^3)], [8, -invep0 * fsckv / n^3],
[9, cpack(0, fmkz * fx3v / n^3)], [10, cpack(0, -fmkz * fx2v / n^3)]],
[3, [1, cpack(0, fmkz * fsckv / n^3)], [2, cpack(0, -fmky * fsckv / n^3)],
[3, fcv / n^3], [7, cpack(0, -fmkz * fx1v / n^3)],
[8, cpack(0, fmky * fx1v / n^3)]],
[4, [0, cpack(0, -fmkz * fsckv / n^3)], [2, cpack(0, fmkx * fsckv / n^3)],
[4, fcv / n^3], [6, cpack(0, fmkz * fx1v / n^3)],
[8, cpack(0, -fmkx * fx1v / n^3)]],
[5, [0, cpack(0, fmky * fsckv / n^3)], [1, cpack(0, -fmkx * fsckv / n^3)],
[5, fcv / n^3], [6, cpack(0, -fmky * fx1v / n^3)],
[7, cpack(0, fmkx * fx1v / n^3)]] ])
```

*rmat = complicated location-dependent convolution kernel  
in frequency domain, computed from multiple seed tables*



# From FFTX DAG to Generated Single Size Code



# Performance Engineering for SYCL, CUDA, and HIP

Transform  
user specified

DFT<sub>8</sub>



Fast algorithm  
in SPL  
many choices

$$(DFT_2 \otimes I_4) T_4^8 (I_2 \otimes ((DFT_2 \otimes I_2) \cdot T_2^4 (I_2 \otimes DFT_2) L_2^4)) L_2^8$$



Σ-SPL:

$$\sum (S_j DFT_2 G_j) \sum \left( \sum (S_{k,l} \text{diag}(t_{k,l}) DFT_2 G_l) \sum (S_m \text{diag}(t_m) DFT_2 G_{k,m}) \right)$$



C Code:

```
void sub(double *y, double *x) {
    double f0, f1, f2, f3, f4, f7, f8, f10, f11;
    f0 = x[0] - x[3];
    f1 = x[0] + x[3];
    f2 = x[1] - x[2];
    f3 = x[1] + x[2];
    f4 = f1 - f3;
    y[0] = f1 + f3;
    y[2] = 0.7071067811865476 * f4;
    f7 = 0.9238795325112867 * f0;
    f8 = 0.3826834323650898 * f2;
    y[1] = f7 + f8;
    f10 = 0.3826834323650898 * f0;
    f11 = (-0.9238795325112867) * f2;
    y[3] = f10 + f11;
}
```

Optimization at all  
abstraction levels



parallelization  
vectorization  
SIMTification



loop and locality  
optimizations



constant folding  
scheduling  
ISA specific  
optimizations  
.....

**MDDFT Pease Algorithm**

$$DFT_{n \times n \times n} \rightarrow \prod_{i=0}^{n-1} (DFT_n \otimes I_{n^2}) L_n^{n^3}$$

**SVCT for SIMT transformation**

$$(A_m \otimes L_m^{mn}) \rightarrow (L_m^{mp} (I_p \otimes A_m) \otimes I_{n/p}) (I_p \otimes L_m^{mn/p})$$

**Loop flattening and regularization**

$$\sum_{i=0}^{m-1} \sum_{j=0}^{n-1} A_{i,j} \rightarrow \sum_{k=0}^{mn-1} A_{\lfloor k/n \rfloor, k \bmod n}$$

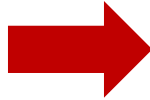
```
simt_loop(i, [0..11], 16, op)
->
simt_loop(i, [0..15], 16,
    if(i <= 11, op, skip()))
```

**ISA specific conf code**

```
cudaDeviceSetLimit(cudaLimitMallocHeapSize,
    1073741824);
cudaFuncSetCacheConfig(ker_exaefl_kernel10,
    cudaFuncCachePreferL1);
cudaFuncSetCacheConfig(ker_exaefl_kernel11,
    cudaFuncCachePreferL1);
cudaMalloc(((void *) &(hp1)),
    (sizeof(double) * 270336));
cudaMemcpyToSymbol(P1, &(hp1),
    sizeof(double *));
```

# How Does SPIRAL JIT Actually Work?

```
// CUDA kernel
__global__ void ker_mdprdf3d_4x4x42(double *Y) {
    double s100, s101, s102, s97, s98, s99;
    int a191, a192;
    a191 = ((24*blockIdx.y) + (6*blockIdx.x));
    s97 = P2[a191];
    s98 = P2[(a191 + 4)];
    s99 = (s97 + s98);
    s100 = (s97 - s98);
    s101 = (2.0*P2[(a191 + 2)]);
    s102 = (2.0*P2[(a191 + 3)]);
    a192 = ((16*blockIdx.y) + (4*blockIdx.x));
    Y[a192] = (s99 + s101);
    Y[(a192 + 2)] = (s99 - s101);
    Y[(a192 + 1)] = (s100 + s102);
    Y[(a192 + 3)] = (s100 - s102);
}
```



```
// main transform function, calls kernels
void mdprdf3d_4x4x4(double *Y, double *X) {
    dim3 b37(1, 1, 1), b38(1, 1, 1), b39(1, 1, 1),
    g1(4, 3, 1), g2(4, 3, 1), g3(4, 4, 1);
    ker_mdprdf3d_4x4x40<<<g1, b37>>>(X);
    ker_mdprdf3d_4x4x41<<<g2, b38>>>();
    ker_mdprdf3d_4x4x42<<<g3, b39>>>(Y);
}
```

```
// kernels as strings for the CUDA JIT
kernels[2] = { "__global__ void ker_mdprdf3d_4x4x42(double *Y) {\n"
    "    double s100, s101, s102, s97, s98, s99;\n"
    "    int a191, a192;\n"
    "    a191 = ((24*blockIdx.y) + (6*blockIdx.x));\n"
    "    s97 = P2[a191];\n"
    "    s98 = P2[(a191 + 4)];\n"
    "    s99 = (s97 + s98);\n"
    "    s100 = (s97 - s98);\n"
    "    s101 = (2.0*P2[(a191 + 2)]);\n"
    "    s102 = (2.0*P2[(a191 + 3)]);\n"
    "    a192 = ((16*blockIdx.y) + (4*blockIdx.x));\n"
    "    Y[a192] = (s99 + s101);\n"
    "    Y[(a192 + 2)] = (s99 - s101);\n"
    "    Y[(a192 + 1)] = (s100 + s102);\n"
    "    Y[(a192 + 3)] = (s100 - s102);\n"
    "} \n", 1, DOUBLE_PY};
```

```
// main transform function encoded as integer table
transform = {MDPRDFT, 3, {4, 4, 4}, 3,
    { kernels[0], dim3(4, 3, 1), dim3(1, 1, 1), 1, DOUBLE_PX },
    { kernels[1], dim3(4, 3, 1), dim3(1, 1, 1), 0 },
    { kernels[2], dim3(4, 4, 1), dim3(1, 1, 1), 1, DOUBLE_PY }};

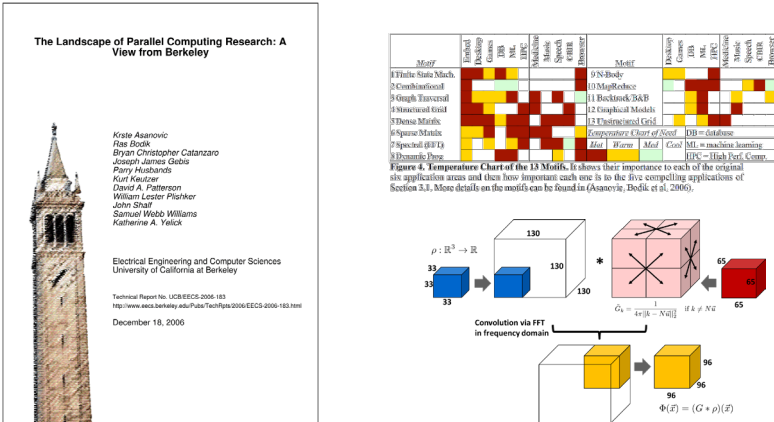
void mdprdf3d_4x4x4(double *Y, double *X) {
    // one-time JITs kernel, interprets and executes integer table
    execute(transform, MODE_JIT_CACHE);
}
```

# Bigger Ecosystem: ECP, X-Stack, and DARPA

## Multi-language, Multi target

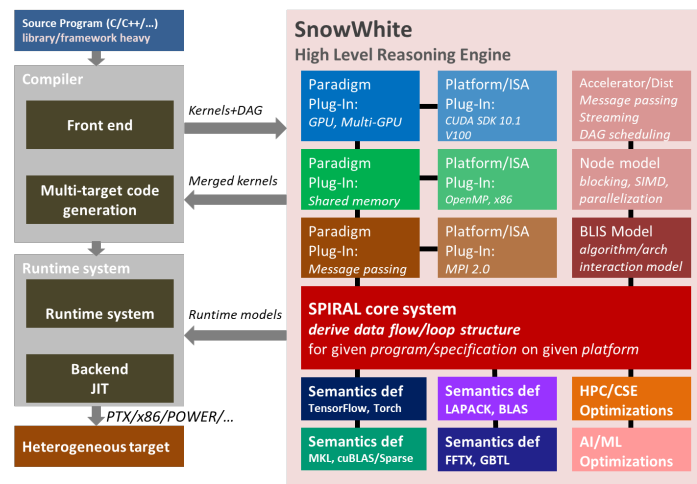


## Cross-motif optimization



Cross-call, cross-library, cross-motif

## SnowWhite: SPIRAL inside compilers



## LibraryX, powered by SPIRAL

